

```
from transformers import pipeline
```

```
text2text_generator = pipeline("text2text-generation")
text2text_generator("question: What is 42 ? context: 42 is the answer to life, the universe and everything")
[{'generated_text': 'the answer to life, the universe and everything'}]
```

```
text2text_generator("translate from English to French: I'm very happy")
[{'generated_text': 'Je suis très heureux'}]
```

⚡ No model was supplied, defaulted to google-t5/t5-base and revision a9723ea (<https://huggingface.co/google-t5/t5-base>). Using a pipeline without specifying a model name and revision in production is not recommended.

/usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning: The secret `HF_TOKEN` does not exist in your Colab secrets. To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>), set it as secret. You will be able to reuse this secret in all of your notebooks. Please note that authentication is recommended but still optional to access public models or datasets.

warnings.warn(

config.json: 100%	1.21k/1.21k [00:00<00:00, 44.6kB/s]
-------------------	-------------------------------------

Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download. For better performance, install the 'hf_xet' package. WARNING:huggingface_hub.file_download:Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP download.

model.safetensors: 100%	892M/892M [00:11<00:00, 25.8MB/s]
-------------------------	-----------------------------------

generation_config.json: 100%	147/147 [00:00<00:00, 10.6kB/s]
------------------------------	---------------------------------

spiece.model: 100%	792k/792k [00:00<00:00, 11.0MB/s]
--------------------	-----------------------------------

tokenizer.json: 100%	1.39M/1.39M [00:00<00:00, 13.7MB/s]
----------------------	-------------------------------------

Device set to use cpu

```
[{'generated_text': 'Je suis très heureux'}]
```

```
from huggingface_hub import login
login(token="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")
```

```
from transformers import pipeline
generator = pipeline('text-generation', model='EleutherAI/gpt-neo-1.3B')
prompt="what is AI"
res=generator(prompt, max_length=150,temperature=0.9)
print(res[0]['generated_text'])
```

⚡ /usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning: The secret `HF_TOKEN` does not exist in your Colab secrets. To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>), set it as secret. You will be able to reuse this secret in all of your notebooks. Please note that authentication is recommended but still optional to access public models or datasets.

warnings.warn(

config.json: 100%	1.35k/1.35k [00:00<00:00, 137kB/s]
-------------------	------------------------------------

model.safetensors: 100%	5.31G/5.31G [00:21<00:00, 255MB/s]
-------------------------	------------------------------------

tokenizer_config.json: 100%	200/200 [00:00<00:00, 28.7kB/s]
-----------------------------	---------------------------------

vocab.json: 100%	798k/798k [00:00<00:00, 12.1MB/s]
------------------	-----------------------------------

merges.txt: 100%	456k/456k [00:00<00:00, 24.5MB/s]
------------------	-----------------------------------

special_tokens_map.json: 100%	90.0/90.0 [00:00<00:00, 13.8kB/s]
-------------------------------	-----------------------------------

Device set to use cpu

Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncation=True` to explicitly truncate inputs. Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.

what is AI, and is it a good idea?

I will start out by stating that I am not claiming AI will kill our existence. AI will, though, do a lot of bad things for humanity as a result of its capabilities.

The bad AI

For most humans it was inevitable that AI would be used in some way to our detriment. We have to consider how AI will ultimately affect our lives and what we can do to mitigate the risks.

```
pip install diffusers
```

⚡ Collecting diffusers

Downloading diffusers-0.32.2-py3-none-any.whl.metadata (18 kB)

Requirement already satisfied: importlib-metadata in /usr/lib/python3/dist-packages (from diffusers) (4.6.4)

Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from diffusers) (3.18.0)

Requirement already satisfied: huggingface-hub<0.23.2 in /usr/local/lib/python3.11/dist-packages (from diffusers) (0.30.1)

Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from diffusers) (2.0.2)

Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.11/dist-packages (from diffusers) (2024.11.6)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from diffusers) (2.32.3)
Requirement already satisfied: safetensors>=0.3.1 in /usr/local/lib/python3.11/dist-packages (from diffusers) (0.5.3)
Requirement already satisfied: Pillow in /usr/local/lib/python3.11/dist-packages (from diffusers) (11.1.0)
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.23.2->diffusers) (2024.10.0)
Requirement already satisfied: packaging>=20.9 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.23.2->diffusers) (24.1)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.23.2->diffusers) (6.0.2)
Requirement already satisfied: tqdm>=4.42.1 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.23.2->diffusers) (4.67.1)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.23.2->diffusers) (4.12.2)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests->diffusers) (3.4.1)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests->diffusers) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests->diffusers) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests->diffusers) (2025.1.31)
Downloading diffusers-0.32.2-py3-none-any.whl (3.2 MB)
3.2/3.2 MB 30.3 MB/s eta 0:00:00
Installing collected packages: diffusers
Successfully installed diffusers-0.32.2

```
import torch
from diffusers import StableDiffusionPipeline, EulerDiscreteScheduler
from PIL import Image


model_id = "stabilityai/stable-diffusion-2"
scheduler = EulerDiscreteScheduler.from_pretrained(model_id, subfolder="scheduler")
pipe = StableDiffusionPipeline.from_pretrained(model_id, scheduler=scheduler, torch_dtype=torch.float16)
pipe = pipe.to("cuda")

prompt = input()
image = pipe(prompt).images[0]
new_image = image.resize((500, 500))
display(new_image)
```

```

/usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it as secret
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
warnings.warn(
scheduler_config.json: 100% 345/345 [00:00<00:00, 37.3kB/s]
model_index.json: 100% 537/537 [00:00<00:00, 46.3kB/s]
Fetching 12 files: 100% 12/12 [00:32<00:00, 3.06s/it]
model.safetensors: 100% 1.36G/1.36G [00:20<00:00, 28.0MB/s]
config.json: 100% 633/633 [00:00<00:00, 53.3kB/s]
vocab.json: 100% 1.06M/1.06M [00:00<00:00, 17.8MB/s]
special_tokens_map.json: 100% 460/460 [00:00<00:00, 12.1kB/s]
merges.txt: 100% 525k/525k [00:00<00:00, 2.56MB/s]
preprocessor_config.json: 100% 342/342 [00:00<00:00, 4.90kB/s]
tokenizer_config.json: 100% 824/824 [00:00<00:00, 9.23kB/s]
config.json: 100% 909/909 [00:00<00:00, 13.3kB/s]
diffusion_pytorch_model.safetensors: 100% 3.46G/3.46G [00:31<00:00, 275MB/s]
diffusion_pytorch_model.safetensors: 100% 335M/335M [00:09<00:00, 12.9MB/s]
config.json: 100% 611/611 [00:00<00:00, 7.56kB/s]
Loading pipeline components...: 100% 6/6 [00:14<00:00, 2.64s/it]
sweet dog
100% 50/50 [00:16<00:00, 3.09it/s]

```



```
pip install -U diffusers
```

```

Collecting diffusers
  Downloading diffusers-0.32.2-py3-none-any.whl.metadata (18 kB)
Requirement already satisfied: importlib-metadata in /usr/lib/python3/dist-packages (from diffusers) (4.6.4)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from diffusers) (3.18.0)
Requirement already satisfied: huggingface-hub>=0.23.2 in /usr/local/lib/python3.11/dist-packages (from diffusers) (0.30.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from diffusers) (2.0.2)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.11/dist-packages (from diffusers) (2024.11.6)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from diffusers) (2.32.3)
Requirement already satisfied: safetensors>=0.3.1 in /usr/local/lib/python3.11/dist-packages (from diffusers) (0.5.3)
Requirement already satisfied: Pillow in /usr/local/lib/python3.11/dist-packages (from diffusers) (11.1.0)
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.23.2->diffusers) (2024.10.0)
Requirement already satisfied: packaging>=20.9 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.23.2->diffusers) (24.1)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.23.2->diffusers) (6.0.2)

```

Requirement already satisfied: tqdm>=4.42.1 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.23.2->diffusers) (4.67.1)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.11/dist-packages (from huggingface-hub>=0.23.2->diffusers) (4.6.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests->diffusers) (3.4.1)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests->diffusers) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests->diffusers) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests->diffusers) (2025.1.31)
Downloading diffusers-0.32.2-py3-none-any.whl (3.2 MB)
3.2/3.2 MB 34.0 MB/s eta 0:00:00
Installing collected packages: diffusers
Successfully installed diffusers-0.32.2

```
from transformers import pipeline
```

```
pipe = pipeline("image-text-to-text", model="llava-hf/llava-interleave-qwen-0.5b-hf")  
messages = [  
    {  
        "role": "user",  
        "content": [  
            {"type": "image", "url": "https://www.ilankelman.org/stopsigns/australia.jpg"},  
            {"type": "text", "text": "What is shown in this image?"},  
        ],  
    },  
]
```

```
outputs = pipe(text=messages, max_new_tokens=60)
```

```
outputs[0]["generated_text"]
```

```
🔄 config.json: 100% 1.25k/1.25k [00:00<00:00, 115kB/s]  
model.safetensors: 100% 1.73G/1.73G [00:06<00:00, 281MB/s]  
Sliding Window Attention is enabled but not implemented for `sdpa`; unexpected results may be encountered.  
generation_config.json: 100% 126/126 [00:00<00:00, 13.8kB/s]  
processor_config.json: 100% 169/169 [00:00<00:00, 19.0kB/s]  
chat_template.json: 100% 699/699 [00:00<00:00, 59.9kB/s]  
preprocessor_config.json: 100% 393/393 [00:00<00:00, 45.2kB/s]  
Using a slow image processor as `use_fast` is unset and a slow processor was saved with this model. `use_fast=True` will be the default  
tokenizer_config.json: 100% 1.60k/1.60k [00:00<00:00, 143kB/s]  
vocab.json: 100% 2.78M/2.78M [00:00<00:00, 6.22MB/s]  
merges.txt: 100% 1.67M/1.67M [00:00<00:00, 3.95MB/s]  
tokenizer.json: 100% 7.03M/7.03M [00:01<00:00, 5.33MB/s]  
added_tokens.json: 100% 101/101 [00:00<00:00, 12.2kB/s]  
special_tokens_map.json: 100% 367/367 [00:00<00:00, 39.9kB/s]  
Device set to use cuda:0
```